



Compiler

Garbage Collection

© Copyright 2005, Matthew B. Dwyer and Robby. The syllabus and lectures for this course are copyrighted materials and may not be used in other course settings outside University of Nebraska-Lincoln and Kansas State University in their current form or modified form without express written permission of one of the copyright holders. During this course, students are prohibited from selling notes to or being paid for taking notes by any person or commercial firm without the express written permission of one of the copyright holders.



Garbage Collector

- is part of the run-time system: it reclaims heap-allocated records that are no longer used.
- A garbage collector should:
 - ☐ reclaim *all* unused records;
 - ☐ spend very little time per record;
 - ☐ not cause significant delays; and
 - ☐ allow all of memory to be used.
- These are difficult and often conflicting requirements.




Without Garbage Collection

- unused records must be explicitly deallocated;
- superior if done correctly;
- but it is easy to miss some records; and
- it is dangerous to handle pointers.



Memory leaks in real life (ical v.2.1)





Which records are in use?


- Ideally, records that will never be accessed in the future execution of the program.
 - but that is of course undecidable...

- Basic conservative assumption:

*A record is **live** if it is reachable from a stack-based program variable.*

- Dead records may still be pointed to by other dead records.

Heap with Live and Dead Records





Mark and Sweep

■ Algorithm

- ☐ explore pointers starting from the program variables, and *mark* all
- ☐ records encountered;
- ☐ *sweep* through all records in the heap and reclaim the unmarked ones; also
- ☐ unmark all marked records.

■ Assumptions:

- ☐ we know the size of each record;
- ☐ we know which fields are pointers; and
- ☐ reclaimed records are kept in a *freelist*.



Mark-and-sweep Code

```
function DFS(x)
  if x is a pointer into the heap then
    if record x is not marked then
      mark record x
    for  $i$  in  $1 \dots |x|$  do
      DFS( $x.f_i$ )
```




Mark-and-sweep Code

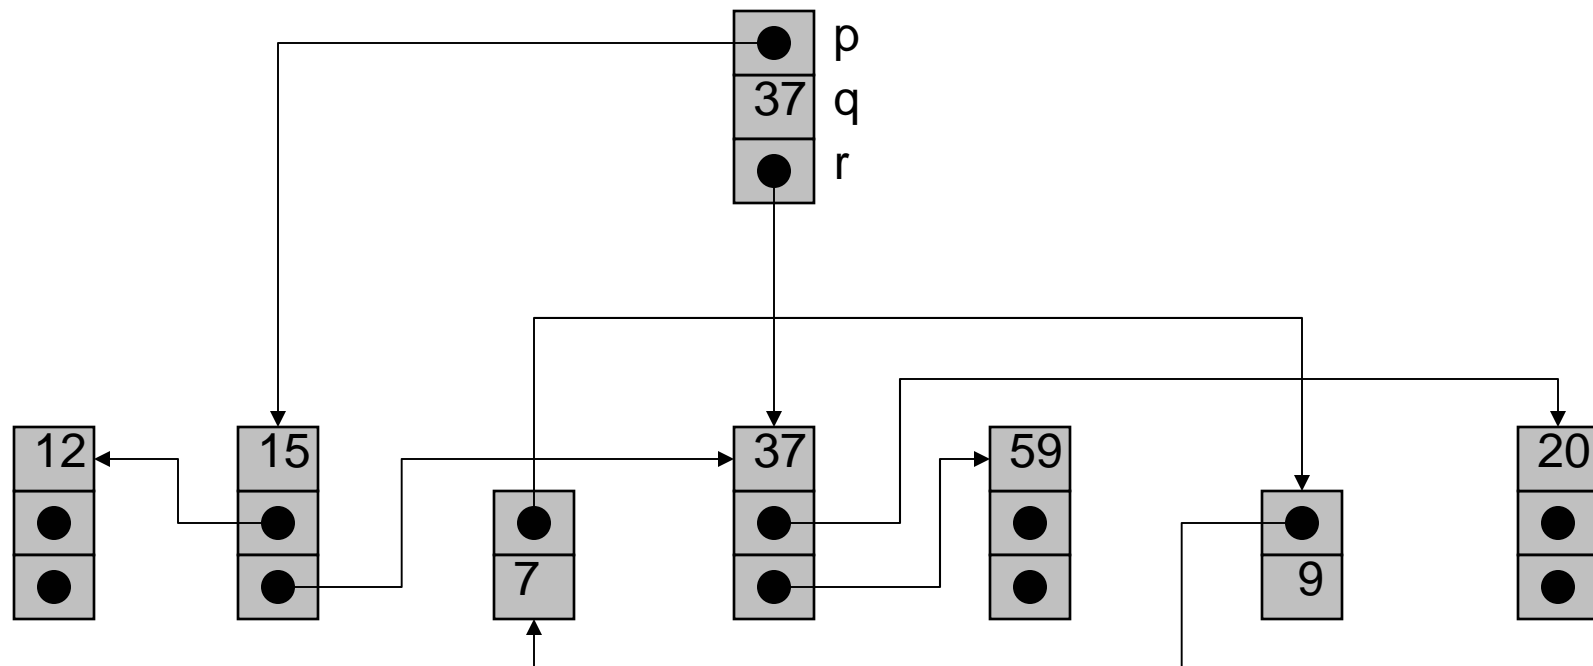
```
function Mark()  
  for each program variable  $v$  do  
    DFS( $v$ )
```



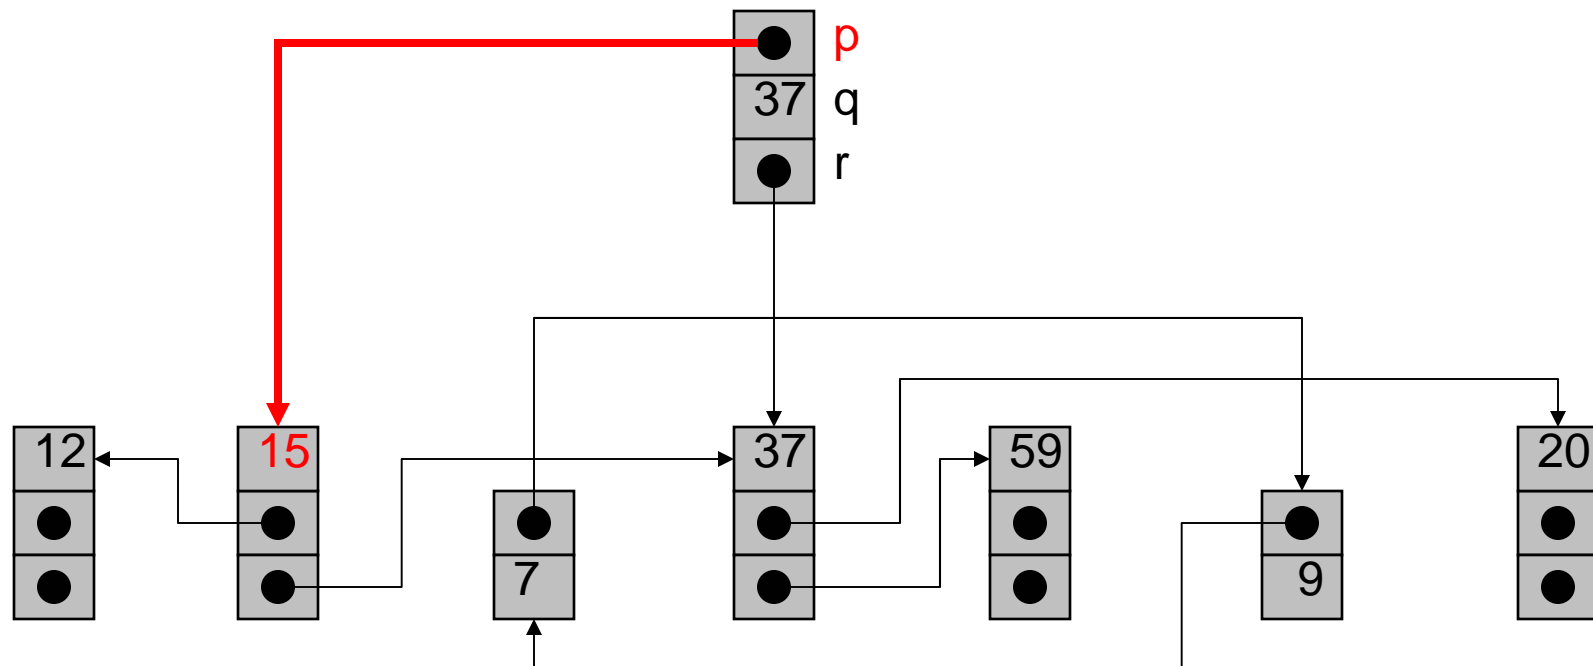
Mark-and-sweep Code

```
function Sweep()  
   $p$  := first address in heap  
  while  $p$  < last address in heap do  
    if record  $p$  is marked then  
      unmark record  $p$   
    else  
       $p.f_1$  := freelist  
      freelist :=  $p$   
       $p$  :=  $p$  + sizeof(record  $p$ )
```

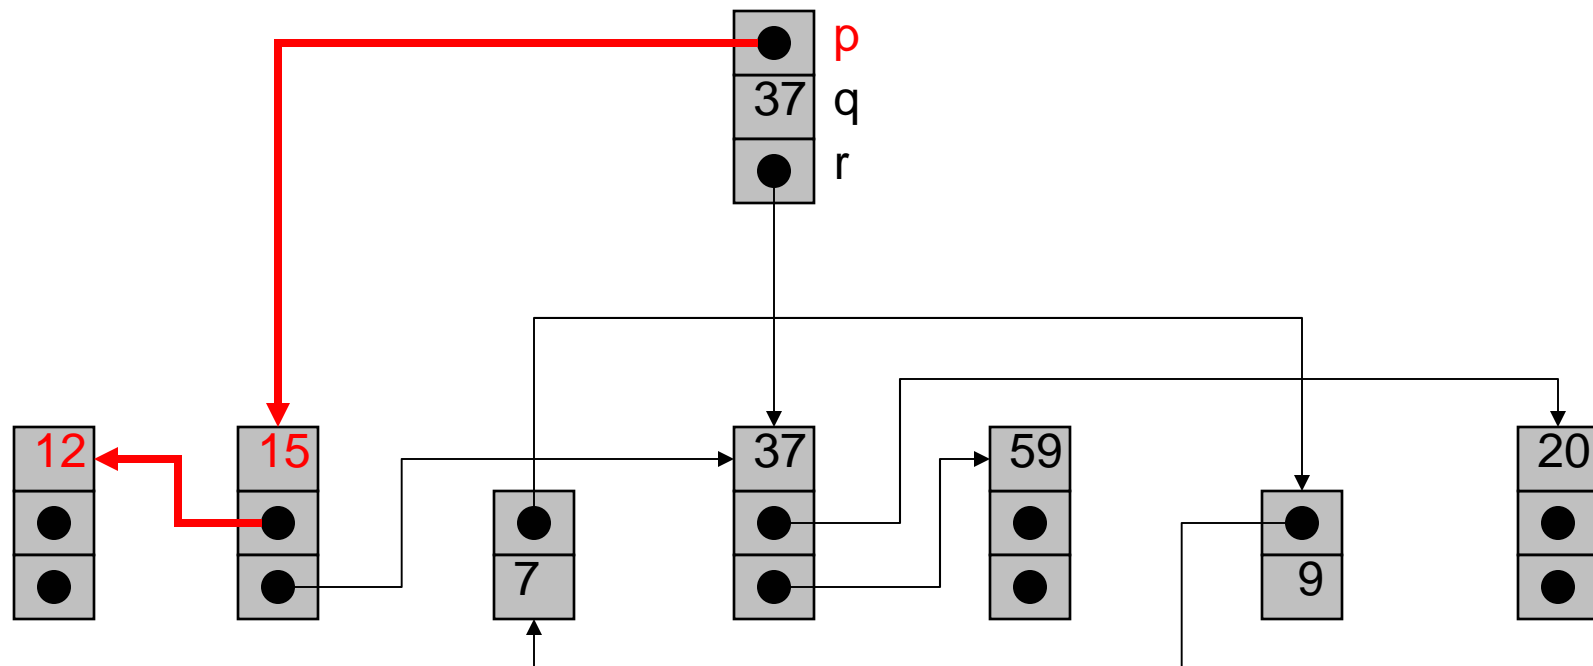
Example



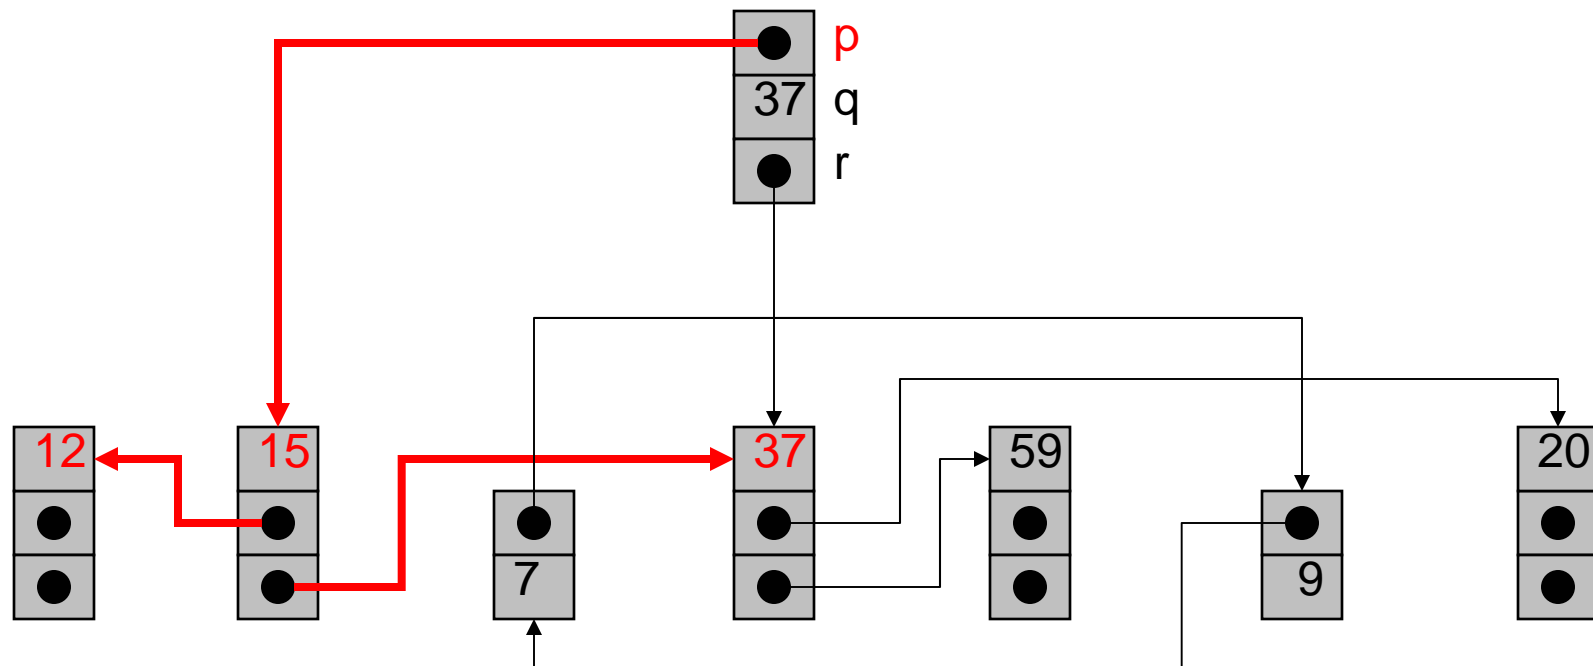
Example



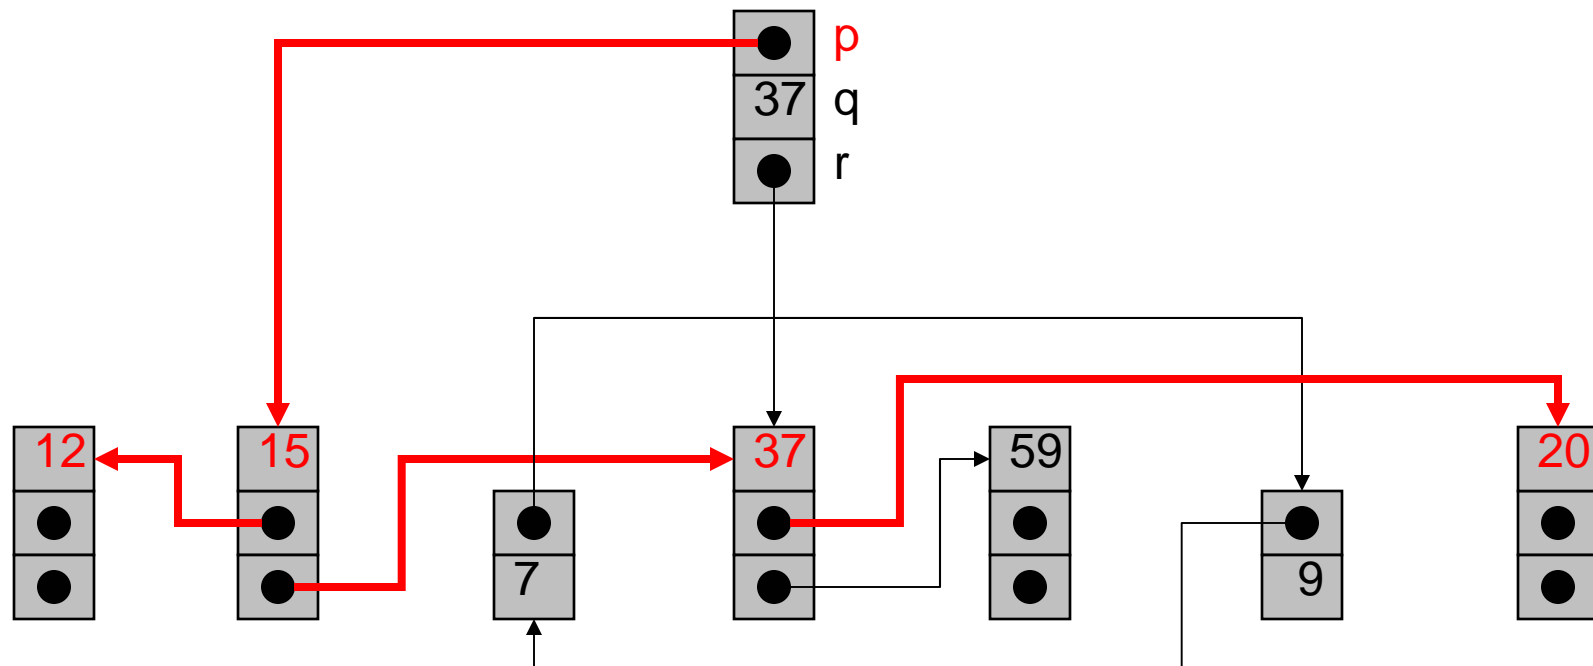
Example



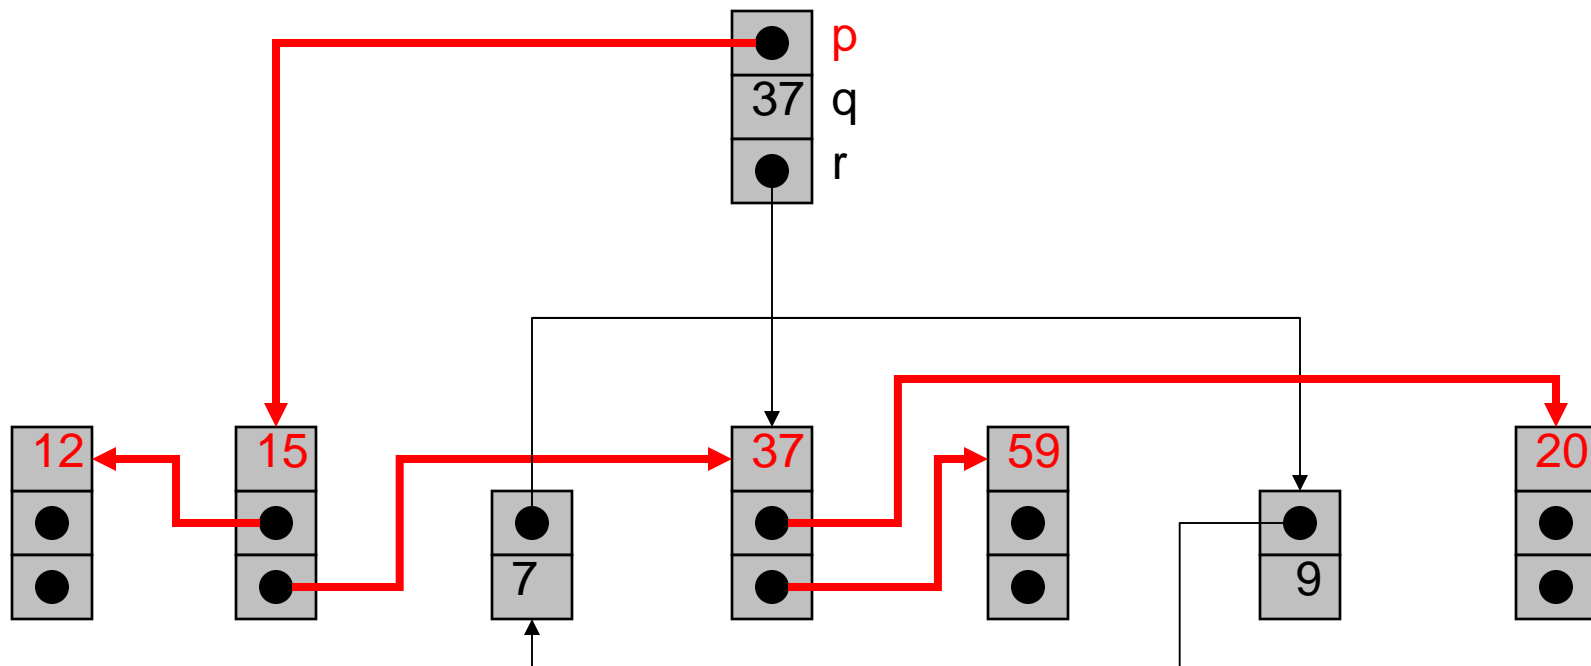
Example



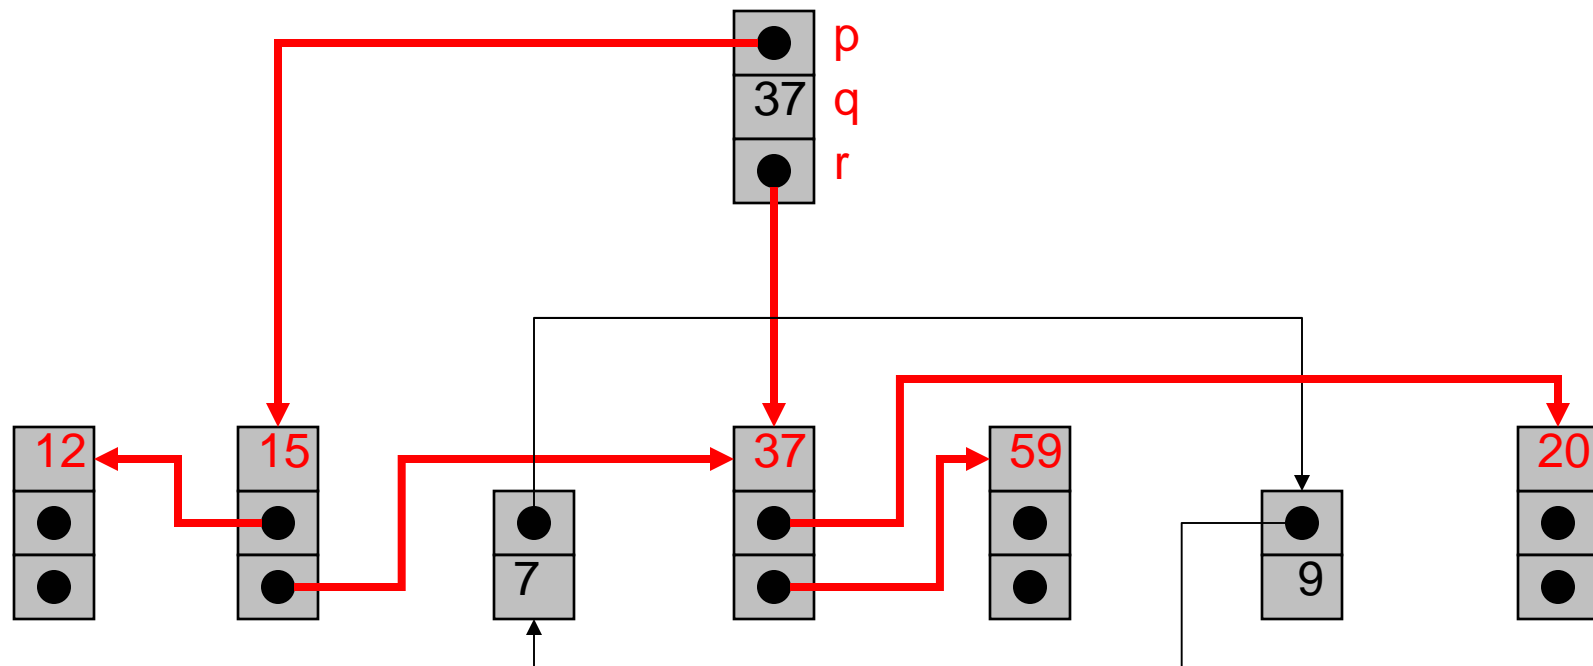
Example




Example



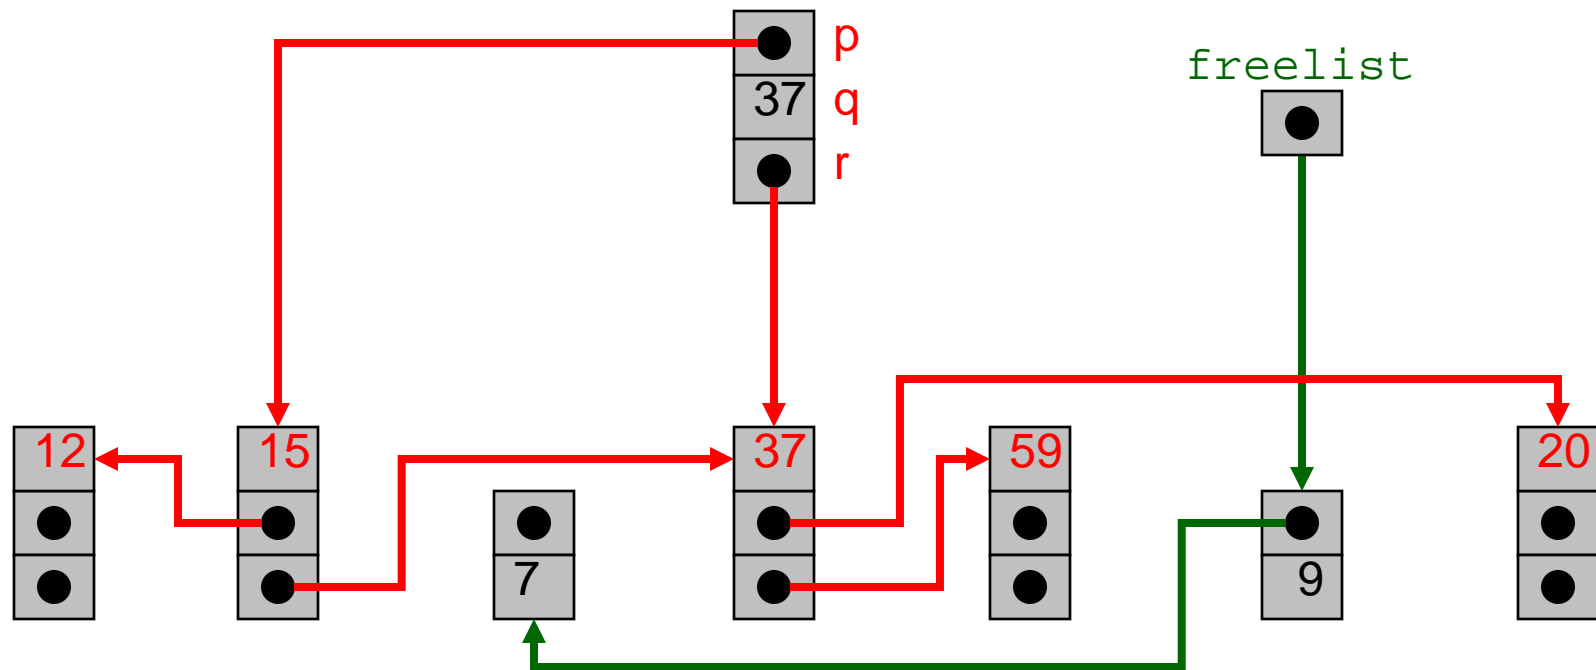
Example



Example



Example





Analysis of Mark-and-Sweep

Assume the heap has R of H words that reachable.

The cost of garbage collection is:

$$c_1R + c_2H$$

Realistic values are:

$$10R + 3H$$

The cost per reclaimed word is:

$$(c_1R + c_2H)(H-R)$$

- if R is close to H , then this is expensive;
- the lower bound is c_2 ;
- increase the heap when $R > 0.5H$; then
- the cost per word is $c_1 + 2c_2 = 16$.



Other Issues

- The DFS recursion stack could have size H (and has at least size $\log H$), which may be too much
 - however, the recursion stack can cleverly be embedded in the fields of marked records (pointer reversal).
- Records can be kept sorted by sizes in the `freelist`. Records may be split into smaller pieces if necessary.
- The heap may become *fragmented*: containing many small free records but none that are large enough.



Reference Counting

■ Algorithm

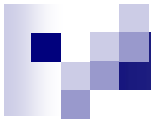
- ☐ maintain a counter of the references to each record;
- ☐ for each assignment, update the counters appropriately; and
- ☐ a record is dead when its counter is zero.

■ Advantages:

- ☐ is simple and attractive;
- ☐ catches dead records immediately; and
- ☐ does not cause long pauses.

■ Disadvantages:

- ☐ cannot detect cycles of dead records; and
- ☐ is much too expensive.



Reference Counting Code

```
function Increment(x)
```

```
    x.count := x.count + 1
```

```
function Decrement(x)
```

```
    x.count := x.count - 1
```

```
    if x.count == 0 then
```

```
        PutOnFreeList(x)
```



Reference Counting Code

```
function PutOnFreelist(x)
```

```
    Decrement( $x.f_1$ )
```

```
     $x.f_1 := \text{freelist}$ 
```

```
    freelist := x
```

```
function RemoveFromFreelist(x)
```

```
    for  $i$  in 2 ...  $|x|$  do
```

```
        Decrement( $x.f_i$ )
```




Stop-and-Copy Counting

■ Algorithm

- divide the heap into two parts;
- only use one part at a time;
- when it runs full, copy live records to the other part; and
- switch the roles of the two parts.


■ Advantages:

- allows fast allocation (no `freelist`);
- avoids fragmentation;
- collects in time proportional to R ; and
- avoids stack and pointer reversal.


■ Disadvantage:

- wastes half your memory.


Example Stop-and-copy




Example Stop-and-copy



Example Stop-and-copy



Example Stop-and-copy





Copy Code

```
function Copy()
  scan := next := start of to-space
  for each program variable  $v$  do
     $v := \text{Forward}(v)$ 
  while scan < next do
    for  $i$  in 1 .. |scan| do
       $\text{scan}.f_i := \text{Forward}(\text{scan}.f_i)$ 
    scan := scan + sizeof(record scan)
```



Copy Code

```
function Forward(p)
  if  $p$  in from-space then
    if  $p.f_1$  in to-space then
      return  $p.f_1$ 
    else
      for  $i$  in 1 ..  $|p|$  do
         $next.f_i := p.f_i$ 
         $p.f_i := next$ 
         $next := next + \text{sizeof}(\text{record } p)$ 
      return  $p.f_1$ 
  else return  $p$ 
```



Analysis of Stop-and-Copy

Assume the heap has R of H words that reachable.
The cost of garbage collection is:

$$c_3R$$

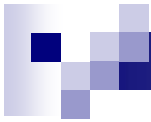
Realistic values are:

$$10R$$

The cost per reclaimed word is:

$$(c_3R)(H/2 - R)$$

- no lower bound as H grows;
- If $H = 4R$ then the cost is constant at approximately 10




Earlier Assumptions

We assumed

- ☐ we know the size of each record; and
- ☐ we know which fields are pointers.

For object-oriented languages, each record already contains a pointer to a class descriptor.

For general languages, we must sacrifice a few bytes per record.



For You To Do

- What algorithm should we use?
- Under what conditions?



Common Algorithms

- Mark-and-sweep or stop-and-copy
- Garbage collection is expensive
 - ~100 instructions for a small object
- Extensions
 - Generational and Region-based Collection
 - Identify objects that die around same time and collect them at once -- faster
 - Incremental and partial collection
 - Reclaim only a part of the heap at a time -- smoother



Generational Collection

- observation: the young die quickly;
- hence the collector should focus on young records;
- divide the heap into generations: G_0, G_1, G_2, \dots ;
- all records in G_i are younger than records in G_{i+1} ;
- collect G_0 often, G_1 less often, and so on; and
- promote a record from G_i to G_{i+1} when it survives several collections.



Generational Collection

■ How to collect the G_0 generation:

- roots are no longer just program variables but also pointers from G_1 , G_2 , ...;
- it might be very expensive to find those pointers;
- fortunately, they are rare; so we can try to remember them.

■ Ways to remember:

- maintain a list of all updated records (use marks to make this a set);
- mark pages of memory that contain updated records (in hardware or software);
- Syntactic extensions (e.g., RTJava scoped memory)



Incremental Collection

Garbage collection may cause long pauses

- this is undesirable for interactive or real-time programs; so
- try to interleave the garbage collection with the program execution.

Two players access the heap:

- the *mutator*. creates records and moves pointers around; and
- the *collector*. tries to collect garbage.

Some invariants are clearly required to make this work.

- The mutator will suffer some slowdown to maintain these invariants.